



SANTANDER PRODUCT RECOMMENDATION

By Sarah Esayas

OBJECTIVE

With 1.5 years of customers' behavior data from Santander Bank, as well as demographic data, we are to predict what kind of product customers would purchase.

Products ranged from payroll accounts to long-term deposits to loans and credit/debit cards.





**FIRST LOOK AT OUR
DATA...**

	fecha_dato	ncodpers	ind_empleado	pais_residencia	sexo	age	fecha_alta	ind_nuevo	antiguedad	indrel	...	ind_hip_fin_ult1
0	2015-01-28	1375586	N	ES	H	35	2015-01-12	0.0	6	1.0	...	0
1	2015-01-28	1050611	N	ES	V	23	2012-08-10	0.0	35	1.0	...	0
2	2015-01-28	1050612	N	ES	V	23	2012-08-10	0.0	35	1.0	...	0
3	2015-01-28	1050613	N	ES	H	22	2012-08-10	0.0	35	1.0	...	0
4	2015-01-28	1050614	N	ES	V	23	2012-08-10	0.0	35	1.0	...	0
...
99994	2015-01-28	890570	N	ES	V	24	2010-09-02	0.0	58	1.0	...	0
99995	2015-01-28	890561	N	ES	H	25	2010-09-02	0.0	58	1.0	...	0
99996	2015-01-28	890559	N	ES	H	48	2010-09-03	0.0	58	1.0	...	0
99997	2015-01-28	890557	N	ES	V	53	2010-09-02	0.0	58	1.0	...	0
99998	2015-01-28	890498	N	ES	V	47	2010-09-02	0.0	58	1.0	...	0

99999 rows × 48 columns

d_pres_fin_ult1	ind_reca_fin_ult1	ind_tjcr_fin_ult1	ind_valo_fin_ult1	ind_viv_fin_ult1	ind_nomina_ult1	ind_nom_pens_ult1	ind_recibo_ult1
0	0	0	0	0	0.0	0.0	0
0	0	0	0	0	0.0	0.0	0
0	0	0	0	0	0.0	0.0	0
0	0	0	0	0	0.0	0.0	0
0	0	0	0	0	0.0	0.0	0
...
0	0	0	0	0	0.0	0.0	1
0	0	0	0	0	0.0	0.0	0
0	0	0	0	0	0.0	0.0	1
0	0	0	0	0	0.0	0.0	0
0	0	0	0	0	0.0	0.0	0

UNDERSTANDING THE DATA

In [2]:

```
df.dtypes
```

In [3]:

```
# Let's check if our values corresponds with the data type
for i in df.columns:
    print(f"{i}\n", df[i].unique(), "\n", f"Values = {df[i].unique().size}\n")
```

Column Name	Description	Data Type
fecha_datos FIXED	Data starts at 01/28/2015	object DATETIME
ncodpers	Customer code	int64
ind_empleado	0 = A = active 1 = B = ex employed 2 = F = subsidiary 3 = N = not employee 4 = P = passive (?)	object
pais_residencia	Customer's Country residence 114 different countries 0 = 'BO' 1 = 'DE' 2 = 'ES' 3 = 'PY'	object
sexo	0 = H = Male 1 = V = Female	object
age FIXED	Age 115 different ages	object FLOAT64
fecha_alta FIXED	The date in which the customer became as the first holder of a contract in the bank	object DATE/TIME
ind_nuevo FIXED	New customer Index. 0 = Old customer 1 = New customer if the customer registered in the last 6 months.	float64 OBJECT
antiguedad	Customer seniority (in months) ??? Does this mean how long they've had it? 249 different numbers	object INT64 ?
indrel FIXED	Indicador de Relación (Relationship Indicator) 1 (First/Primary) 99 (Primary customer during the month but not at the end of the month)	float64 OBJECT
ult_fec_cli_1t FIXED	Last date as primary customer (if he isn't at the end of the month)	object DATETIME
indrel_1mes FIXED	Customer type at the beginning of the month, 1 (First/Primary customer), 2 (co-owner),	float64 OBJECT

	I = I (inactive), P (former customer), R (Potential)	
indresi	Residence index 1 = S (Yes) or 0 = N (No) if the residence country is the same than the bank country)	object
indext	Foreigner index 1 = (S (Yes) or 0 = N (No) if the customer's birth country is different than the bank country)	object
conyuemp	Spouse index: N-No S-Yes if the customer is spouse of an employee	object
canal_entrada	channel used by the customer to join There are 157 different categories	object
indfall	Deceased index. N/S 0 = N 1 = S	object
tipodom	Address type. 1, primary address Everyone put a primary address	float64 OBJECT
cod_prov	Province code (customer's address) 53 different province codes	float64
nomprov	Province name 53 different province names	object
ind_actividad_cliente FIXED	Activity index 1, active customer 0, inactive customer	float64 OBJECT
renta	Gross income of the household mean = \$139,646.2 sd = 2.389858e+05 min = 1.202730e+03 max = 2.889440e+07	float64
segmento	segmentation: 0 - VIP, 1 - Individuals / Particulares 2 - college graduated nan	object

ind_ahor_fin_ult1	Saving Account	object
ind_aval_fin_ult1	Guarantees	object
ind_cco_fin_ult1	Current Accounts	object
ind_cder_fin_ult1	Derivada Account	object
ind_cno_fin_ult1	Payroll Account	object
ind_ctju_fin_ult1	Junior Account	object
ind_ctma_fin_ult1	Más particular Account	object
ind_ctop_fin_ult1	particular Account	object
ind_ctpp_fin_ult1	particular Plus Account	object
ind_deco_fin_ult1	Short-term deposits	object
ind_deme_fin_ult1	Medium-term deposits	object
ind_dela_fin_ult1	Long-term deposits	object
ind_ecue_fin_ult1	e-account	object
ind_fond_fin_ult1	Funds	object
ind_hip_fin_ult1	Mortgage	object
ind_plan_fin_ult1	Pensions	object
ind_pres_fin_ult1	Loans	object
ind_reca_fin_ult1	Taxes	object
ind_tjcr_fin_ult1	Credit Card	object
ind_valo_fin_ult1	Securities	object
ind_viv_fin_ult1	Home Account	object
ind_nomina_ult1	Payroll	object
	** nan values	
ind_nom_pens_ult1	Pensions	object
	** nan values	
ind_recibo_ult1	Direct Debit	object



ON TO DATA PREPROCESSING!

FIRST STEP, CHANGING DATA TYPES:

In [5]:

```
# CHANGING DATA TYPES
```

```
# Dates
```

```
df["fecha_alta"]=pd.to_datetime(df["fecha_alta"])
```

```
df["ult_fec_cli_1t"]=pd.to_datetime(df["ult_fec_cli_1t"])
```

```
# Numeric
```

```
df.loc[df['antiguedad']==' NA','antiguedad']=None
```

```
df["antiguedad"]=pd.to_numeric(df["antiguedad"])
```


```
df.loc[df['age']==' NA','age']=None
```

```
df["age"]=pd.to_numeric(df["age"])
```


NEXT - MISSING DATA!

```
In [6]: # MISSING DATA
missing_data=df.isnull().sum()
print("Missing data:\n", missing_data)
```

```
Missing data:
 ind_employed      683
 pais_residencia   683
 sexo              683
 age              683
 fecha_alta        683
 ind_nuevo         683
 antigüedad        683
 indrel            683
 ult_fec_cli_1t    99871
 indrel_1mes       683
 tiprel_1mes       683
 indresi           683
 indext            683
 canal_entrada     688
 indfall           683
 cod_prov          769
 nomprov           769
 ind_actividad_cliente 683
 renta            18283
 segmento          691
 ind_cco_fin_ult1  0
 ind_cder_fin_ult1 0
 ind_cno_fin_ult1  0
 ind_ctju_fin_ult1 0
 ind_ctma_fin_ult1 0
 ind_ctop_fin_ult1 0
 ind_ctpp_fin_ult1 0
 ind_deco_fin_ult1 0
 ind_deme_fin_ult1 0
 ind_dela_fin_ult1 0
 ind_ecue_fin_ult1 0
 ind_fond_fin_ult1 0
 ind_hip_fin_ult1  0
 ind_plan_fin_ult1 0
 ind_pres_fin_ult1 0
 ind_reca_fin_ult1 0
 ind_tjcr_fin_ult1 0
 ind_valo_fin_ult1 0
 ind_viv_fin_ult1  0
 ind_nomina_ult1   210
 ind_nom_pens_ult1 210
 ind_recibo_ult1   0
 dtype: int64
```



```
In [8]: df.drop('ult_fec_cli_1t', axis=1, inplace=True)
```

```
In [9]: # Let's removing the rows with missing values

missing = ['ind_empleado', 'pais_residencia', 'sexo',
           'age', 'fecha_alta', 'ind_nuevo', 'antiguedad', 'indrel',
           'indrel_1mes', 'tiprel_1mes', 'indresi', 'indext',
           'canal_entrada', 'indfall', 'cod_prov',
           'nomprov', 'ind_actividad_cliente', 'segmento', 'ind_nomina_ult1', 'ind_nom_pens_ult1']

df = df.dropna(subset=missing)
```



Did not delete missing rows of 'renta', instead used implementation with median:

```
In [11]: # Let's replace renta's missing with its median

median_renta = df['renta'].median()

df['renta'].fillna(median_renta, inplace=True)
```



DEALING WITH DUPLICATES:

```
In [12]: duplicates=df.duplicated().sum()  
         print("Number of Duplicate Records: ", duplicates)
```

Number of Duplicate Records: 2722

```
In [13]: df=df.drop_duplicates()
```

```
In [14]: duplicates=df.duplicated().sum()  
         print("Number of Duplicate Records: ", duplicates)
```

Number of Duplicate Records: 0

INVALID ENTRIES AND OUTLIERS:

In [15]:


```
features = ['ind_empleado', 'pais_residencia', 'sexo', 'age', 'fecha_alta',  
            'ind_nuevo', 'antiguedad', 'indrel', 'indrel_1mes', 'tiprel_1mes',  
            'indresi', 'indext', 'canal_entrada', 'indfall', 'cod_prov', 'nomprov',  
            'ind_actividad_cliente', 'renta', 'segmento']  
  
df[features].describe()
```

[15]:	age	fecha_alta	ind_nuevo	antiguedad	indrel	indrel_1mes	cod_prov	ind_actividad_cliente	renta
count	96495.000000	96495	96495.000000	96495.000000	96495.000000	96495.0	96495.000000	96495.000000	9.649500e+04
mean	29.492481	2012-06-30 21:50:33.446292736	0.000114	36.378776	1.126950	1.0	25.039038	0.417027	1.115668e+05
min	4.000000	2002-06-14 00:00:00	0.000000	1.000000	1.000000	1.0	1.000000	0.000000	2.539800e+03
25%	23.000000	2012-06-18 00:00:00	0.000000	33.000000	1.000000	1.0	11.000000	0.000000	6.683498e+04
50%	25.000000	2012-09-06 00:00:00	0.000000	34.000000	1.000000	1.0	28.000000	0.000000	8.961030e+04
75%	31.000000	2012-10-25 00:00:00	0.000000	37.000000	1.000000	1.0	36.000000	1.000000	1.222534e+05
max	106.000000	2015-01-26 00:00:00	1.000000	157.000000	99.000000	1.0	52.000000	1.000000	2.425324e+07
std	10.649402	NaN	0.010676	6.602483	3.524922	0.0	13.656111	0.493070	1.469700e+05


Indrel_1mes had only 1 value after cleaning, so I dropped this column

OUTLIERS


Histogram of age



Histogram of antigüedad



Histogram of renta



```
[18]: # Checking for outliers

def detect_outliers(column):
    Q1 = column.quantile(0.25)
    Q3 = column.quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - (1.5 * IQR)
    upper_bound = Q3 + (1.5 * IQR)
    return (column < lower_bound) | (column > upper_bound)

indep_n = ['age', 'antigüedad', 'renta']

for col in indep_n:
    outliers = detect_outliers(df[col])
    num_outliers = outliers.sum()
    total_values = len(df[col])
    percentage_outliers = (num_outliers / total_values) * 100
    print(f'{col} Outliers: {num_outliers} ({percentage_outliers:.2f}%)\n')
```

age Outliers: 10387 (10.76%)

antigüedad Outliers: 21906 (22.70%)

renta Outliers: 7638 (7.92%)

OUTLIERS

```
[25]: def handle_outliers(col):  
    Q1 = df[col].quantile(0.25)  
    Q3 = df[col].quantile(0.75)  
    IQR = Q3 - Q1  
    lower_bound = Q1 - (1.5 * IQR)  
    upper_bound = Q3 + (1.5 * IQR)  
  
    df.loc[df[col]>upper_bound, col] = Q3  
    df.loc[df[col]<lower_bound, col] = Q1  
  
    for col in indep_n:  
        handle_outliers(col)  
  
    for col in indep_n:  
        d = detect_outliers(df[col]).sum()  
        print(f"{col} Outliers: {d}\n")
```

age Outliers: 0


antigüedad Outliers: 0

renta Outliers: 0



DATA VISUALIZATION

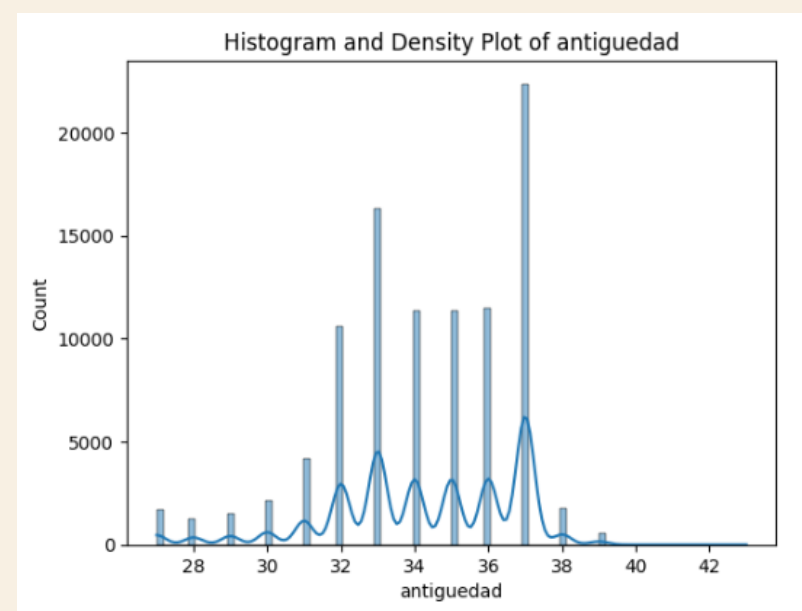

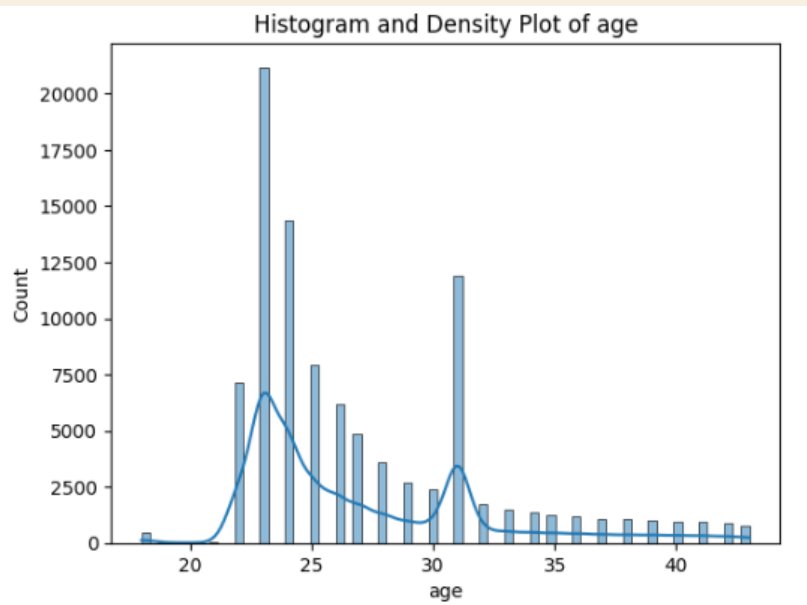
CHECKING TO SEE HOW MANY PRODUCTS WERE BOUGHT:



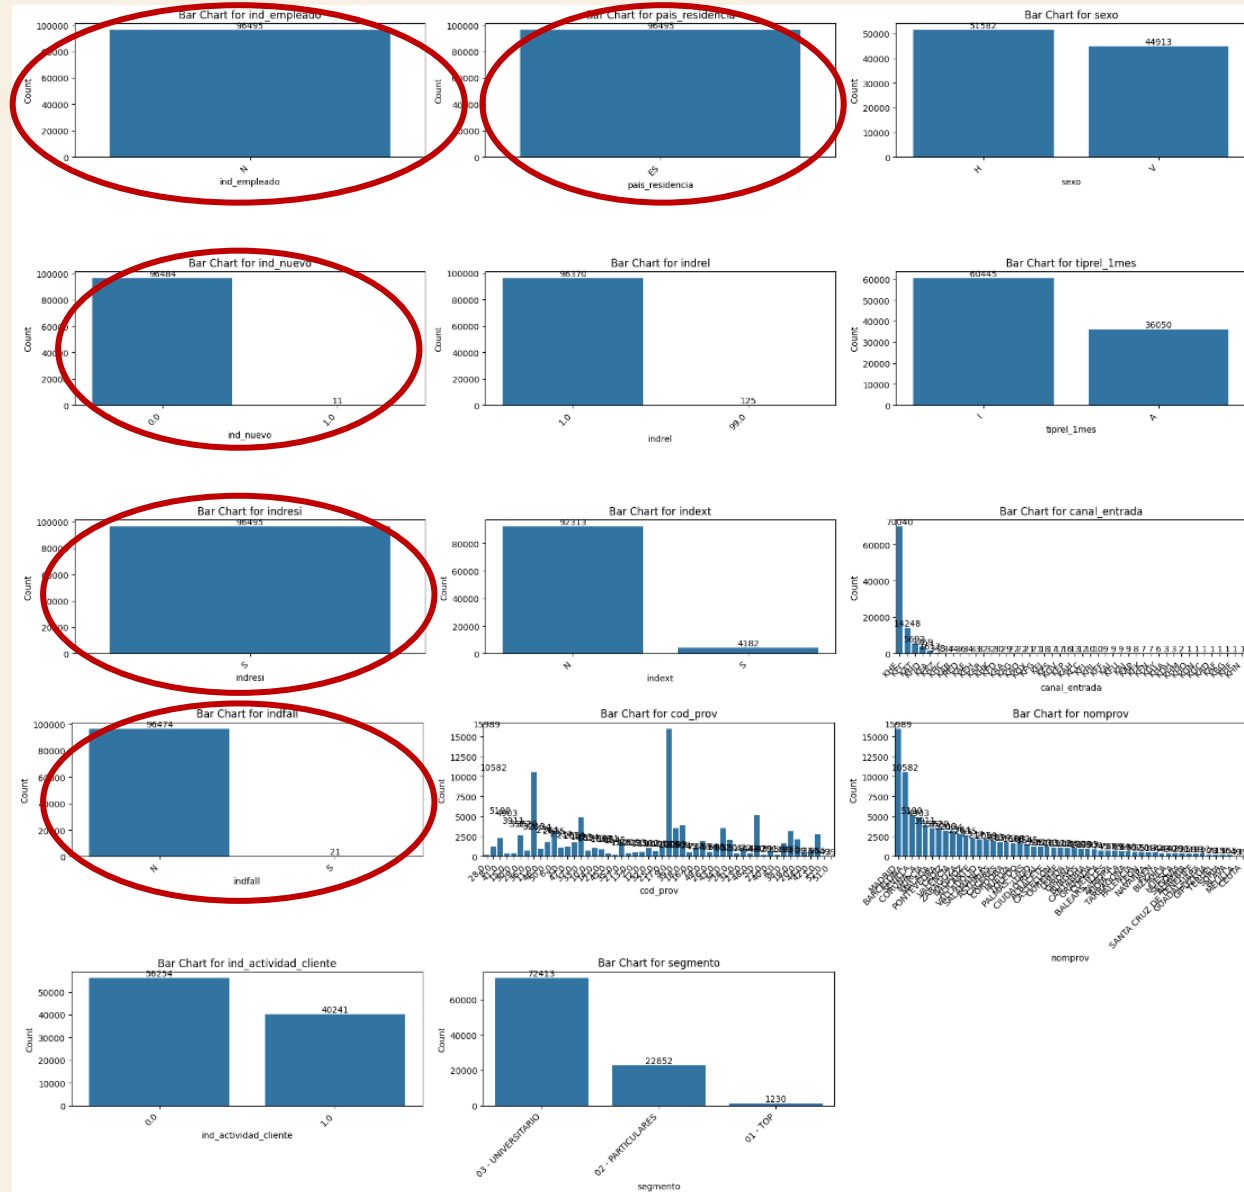
```
ind_cco_fin_ult1      87193.0
ind_recibo_ult1       9695.0
ind_ecue_fin_ult1     7053.0
ind_cno_fin_ult1      5355.0
ind_nom_pens_ult1     3522.0
ind_nomina_ult1       3268.0
ind_dela_fin_ult1     2087.0
ind_tjcr_fin_ult1     1851.0
ind_reca_fin_ult1     1827.0
ind_ctma_fin_ult1     1576.0
ind_ctju_fin_ult1      925.0
ind_valo_fin_ult1     456.0
ind_fond_fin_ult1     370.0
ind_plan_fin_ult1     126.0
ind_deme_fin_ult1      34.0
ind_pres_fin_ult1     19.0
ind_deco_fin_ult1     18.0
ind_hip_fin_ult1      11.0
ind_ctpp_fin_ult1     10.0
ind_ctop_fin_ult1      6.0
ind_viv_fin_ult1       6.0
ind_cder_fin_ult1      4.0
dtype: float64
```

Then chose to focus on products that had > 1000 counts.

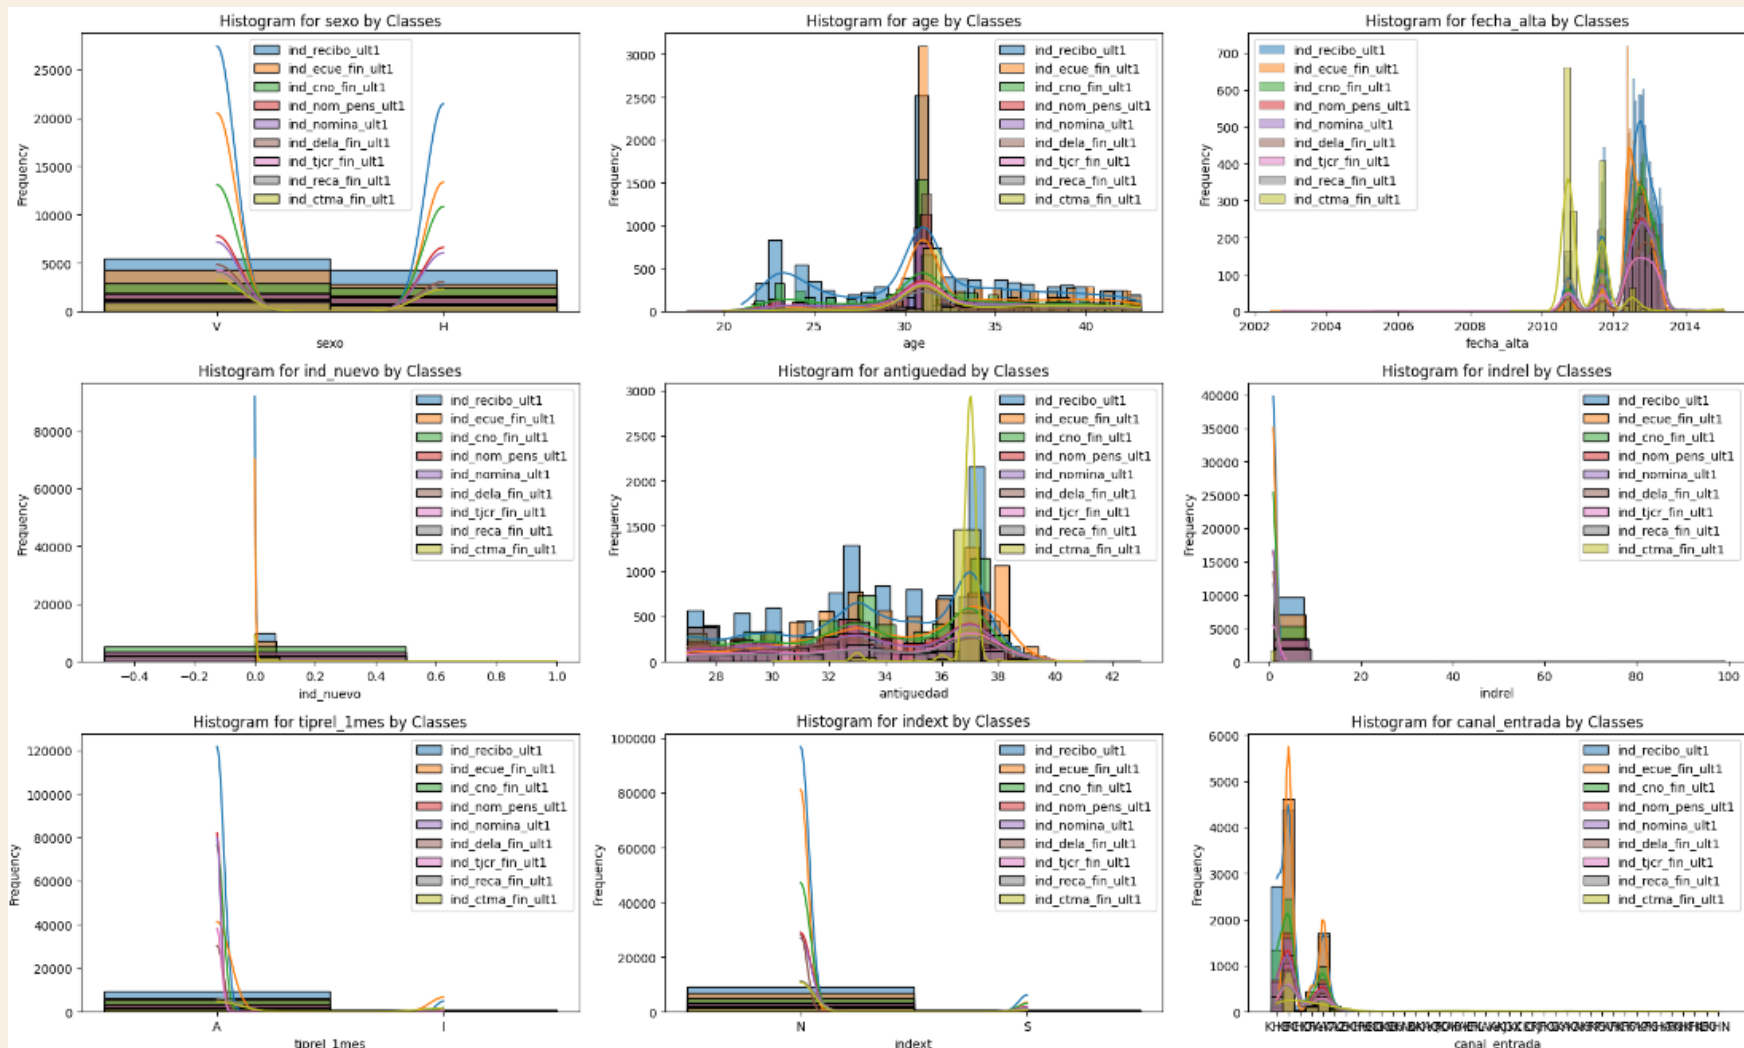
CHECKING FREQUENCY OF NUMERICAL COLUMNS:

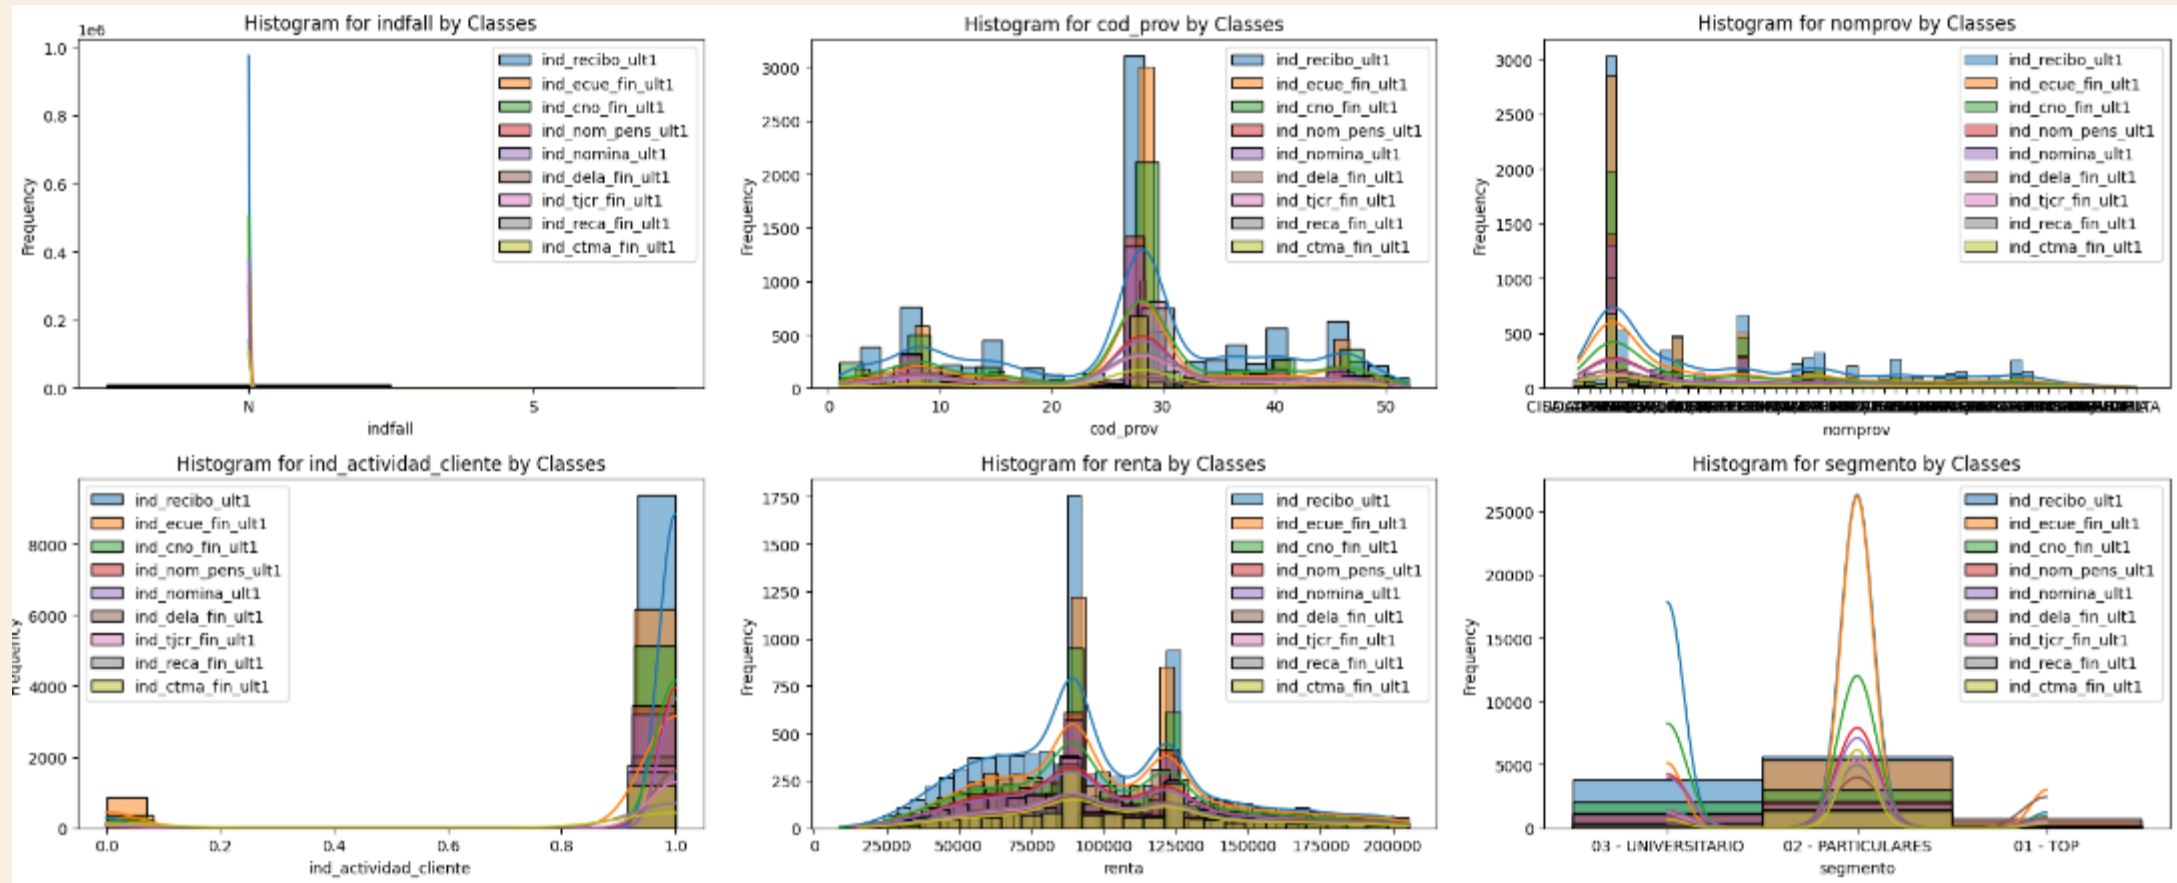


CATEGORICAL COLUMNS:




VISUALIZING RELATIONSHIPS






CORRELATION WITH NUMERICAL VARIABLES



CORRELATION WITH ENCODED CATEGORICAL VARIABLES





MACHINE LEARNING



SELECTED FEATURES AND TARGET VARIABLES

- `'sexo_encoded'` = Gender of the customer (categorical: 0 for male, 1 for female)
- `'age'` = age of customer
- `'antiguedad'` = seniority/tenure of customer in months
- `'tiprel_1mes'` = customer relation type at beginning of month (0 = active, 1 = inactive)
- `'cod_prov'` = province code of customer's residence
- `'ind_actividad_cliente'` = indicator if customer is active or not (0 = ACTIVE, 1 = inactive)
- `'renta'` = gross income of household the customer resides



SELECTED FEATURES AND TARGET VARIABLES

- *'ind_cco_fin_ult1'* = Current accounts
- *'ind_recibo_ult1'* = Direct debits
- *'ind_ecue_fin_ult1'* = e-account
- *'ind_cno_fin_ult1'* = Payroll account
- *'ind_nom_pens_ult1'* = Pensions
- *'ind_nomina_ult1'* = Payroll
- *'ind_dela_fin_ult1'* = Taxes
- *'ind_tjcr_fin_ult1'* = Credit Cards
- *'ind_reca_fin_ult1'* = Taxes
- *'ind_ctma_fin_ult1'* = Long-term deposits



PERFORMANCE METRICS

I'll be using evaluation metrics suitable for multi-label classification tasks, such as:

- Hamming loss
- Jaccard similarity score
- Precision, recall, and F1-score for each product
- Micro/macro averaged metrics



MODEL SELECTION

I've decided to try two different models for this task:

- **K-Nearest Neighbors (KNN) Model**

- I chose this due to its ease of interpretation in robustness to outliers.
- KNN works by classifying instances based on the majority class among their k nearest neighbors

- **Random Forest Model**

- I chose this for its ability to handle outliers and complex relationships in the data.
- RF builds multiple decision trees and combines their predictions through voting to make the final prediction.

KNN

```
[111]: X = df[features]
       y = df[new_targets]

       # Splitting the dataset
       X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2, random_state=470)

[112]: # Preprocessing for numerical variables

       numerical_features = ['age', 'antiguedad', 'renta']

       numerical_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='median')),
                                              ('scaler', StandardScaler())])

[113]: # Preprocessing for categorical variables

       categorical_features = ['sexo_encoded', 'tiprel_1mes_encoded', 'cod_prov', 'ind_actividad_cliente']

       categorical_transformer = Pipeline(steps = [('imputer', SimpleImputer(strategy='most_frequent')),
                                                  ('onehot', OneHotEncoder(handle_unknown='ignore'))])

[114]: # Combining Preprocessing

       preprocessor = ColumnTransformer(transformers=[('num', numerical_transformer, numerical_features),
                                                    ('cat', categorical_transformer, categorical_features)])

[115]: knn_pipeline=Pipeline(steps=[('preprocessor', preprocessor),
                                   ('classifier', KNeighborsClassifier(n_neighbors=4))])

[116]: # Train
       knn_pipeline.fit(X_train,y_train)
```

Classification Report For KNN:

	precision	recall	f1-score	support
0	0.92	0.93	0.93	17367
1	0.37	0.12	0.18	1944
2	0.46	0.15	0.22	1426
3	0.20	0.03	0.05	1100
4	0.19	0.01	0.03	711
5	0.18	0.01	0.03	669
6	0.35	0.04	0.07	447
7	0.12	0.00	0.01	391
8	0.14	0.00	0.01	360
9	0.21	0.06	0.09	288
micro avg	0.88	0.68	0.77	24703
macro avg	0.31	0.14	0.16	24703
weighted avg	0.74	0.68	0.69	24703
samples avg	0.84	0.81	0.81	24703

Accuracy Score: 0.7577076532462822

RANDOM FOREST

```
[93]: from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
df['sexo_encoded'] = label_encoder.fit_transform(df['sexo'])

[94]: label_encoder = LabelEncoder()
df['tiprel_1mes_encoded'] = label_encoder.fit_transform(df['tiprel_1mes'])

[96]: features = ['sexo_encoded', 'age', 'antiguedad', 'tiprel_1mes_encoded', 'cod_prov',
               'ind_actividad_cliente', 'renta']

[97]: from sklearn.ensemble import RandomForestClassifier

X = df[features]
Y = df[new_targets]

clf = RandomForestClassifier(n_estimators=10)
clf = clf.fit(X, Y)

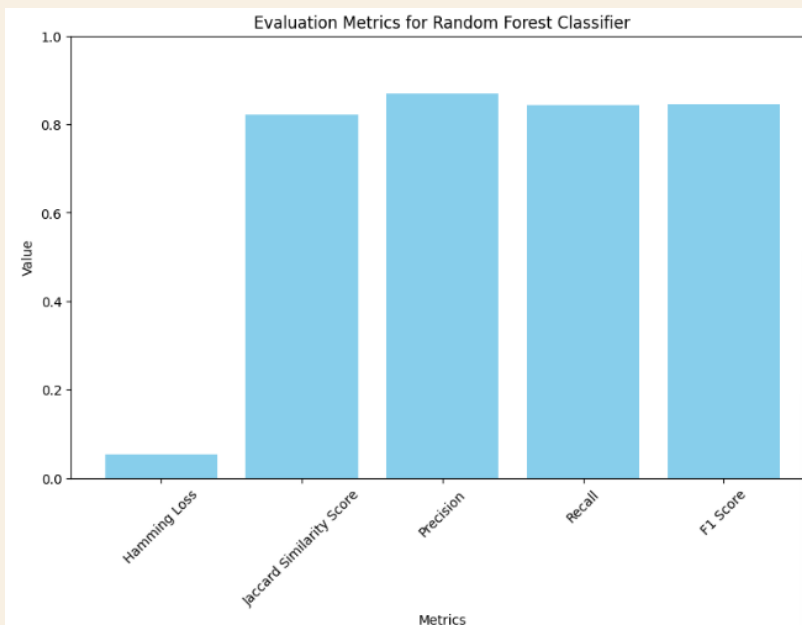
[98]: # Split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

[99]: # Initialize
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

[100]: rf_classifier.fit(X_train, y_train)
```

Accuracy: 0.7688481268459506

Hamming Loss: 0.05405461422871651
Jaccard Similarity Score: 0.8225817276493478
Precision: 0.8698946281450557
Recall: 0.843348779482776
F1 Score: 0.8448592325166241





CONCLUSION

Overall, the Random Forest model demonstrates stronger predictive performance and is better suited for predicting customer product purchases.

In summary, the Random Forest model shows promise in accurately predicting customer product purchases, and efforts should be focused on refining this model for deployment in real-world applications.